



Triton-San: Toward Precise Debugging of Triton Kernels via LLVM Sanitizers



Lechen Yu¹, Tim Lu², Brandon Myers¹, Ofer Dekel¹

¹{yulechen, brandonmyers, oferd}@microsoft.com, ²timthlu@gmail.com

1. Motivation

- Programmers may introduce programming errors in Triton kernels, which may result in
 - Incorrect result
 - Runtime error on GPU
- Multiple kinds of bugs may arise, including:
 - Buffer overflow
 - Use of uninitialized memory
 - Data race
- Silent and non-deterministic bugs could make debugging and reproducing difficult
- Triton doc website provides some guidance for debugging, but it lacks detailed examples and best practices for handling real-world bugs [1].

2. Case Study

Data Races in a Triton Kernel

```
1 @triton.jit
2 def kernel(output_ptr, n, BLOCK_SIZE):
3     pid = tl.program_id(axis=0)
4
5     # Root cause of data race: incorrect block_start
6     # Bug fix: block_start = pid * BLOCK_SIZE
7     block_start = 0
8
9     offsets = block_start + tl.arange(0, BLOCK_SIZE)
10    mask = offsets < n
11    output = tl.full((BLOCK_SIZE, ), 1, dtype=tl.float16)
12
13    # Data races will occur between any two program ids
14    tl.store(output_ptr + offsets, output, mask=mask)
15
16    size = 256
17    output = torch.empty((size, )).to('gpu')
18    grid = lambda meta: (triton.cdiv(size, meta['BLOCK_SIZE']), )
19    kernel[grid](output, size, BLOCK_SIZE=2)
```

3. Triton-San [2]

- We developed a correctness tool tailored for Triton programs
 - currently capable of detecting **buffer overflows** and **data races** in Triton kernels.
- Execute the kernel on CPU through the hardware-agnostic middle-end *triton-shared* [3]
 - Put the focus on programming errors not compiler errors
 - Compile Triton kernels through the built-in **MLIR->LLVM IR->Assembly lowering**
 - Doesn't require a GPU/NPU backend when debugging
- Parallelize kernel execution on CPU via **LLVM OpenMP runtime** [4]
- Leverage **LLVM sanitizers** [5] to ensure reliability and maintainability
 - LLVM sanitizers (from the compiler-rt subproject) have been extensively optimized to detect bugs with high efficiency
 - AddressSanitizer (ASan): catches buffer overflows
 - ThreadSanitizer (TSan): catches data races

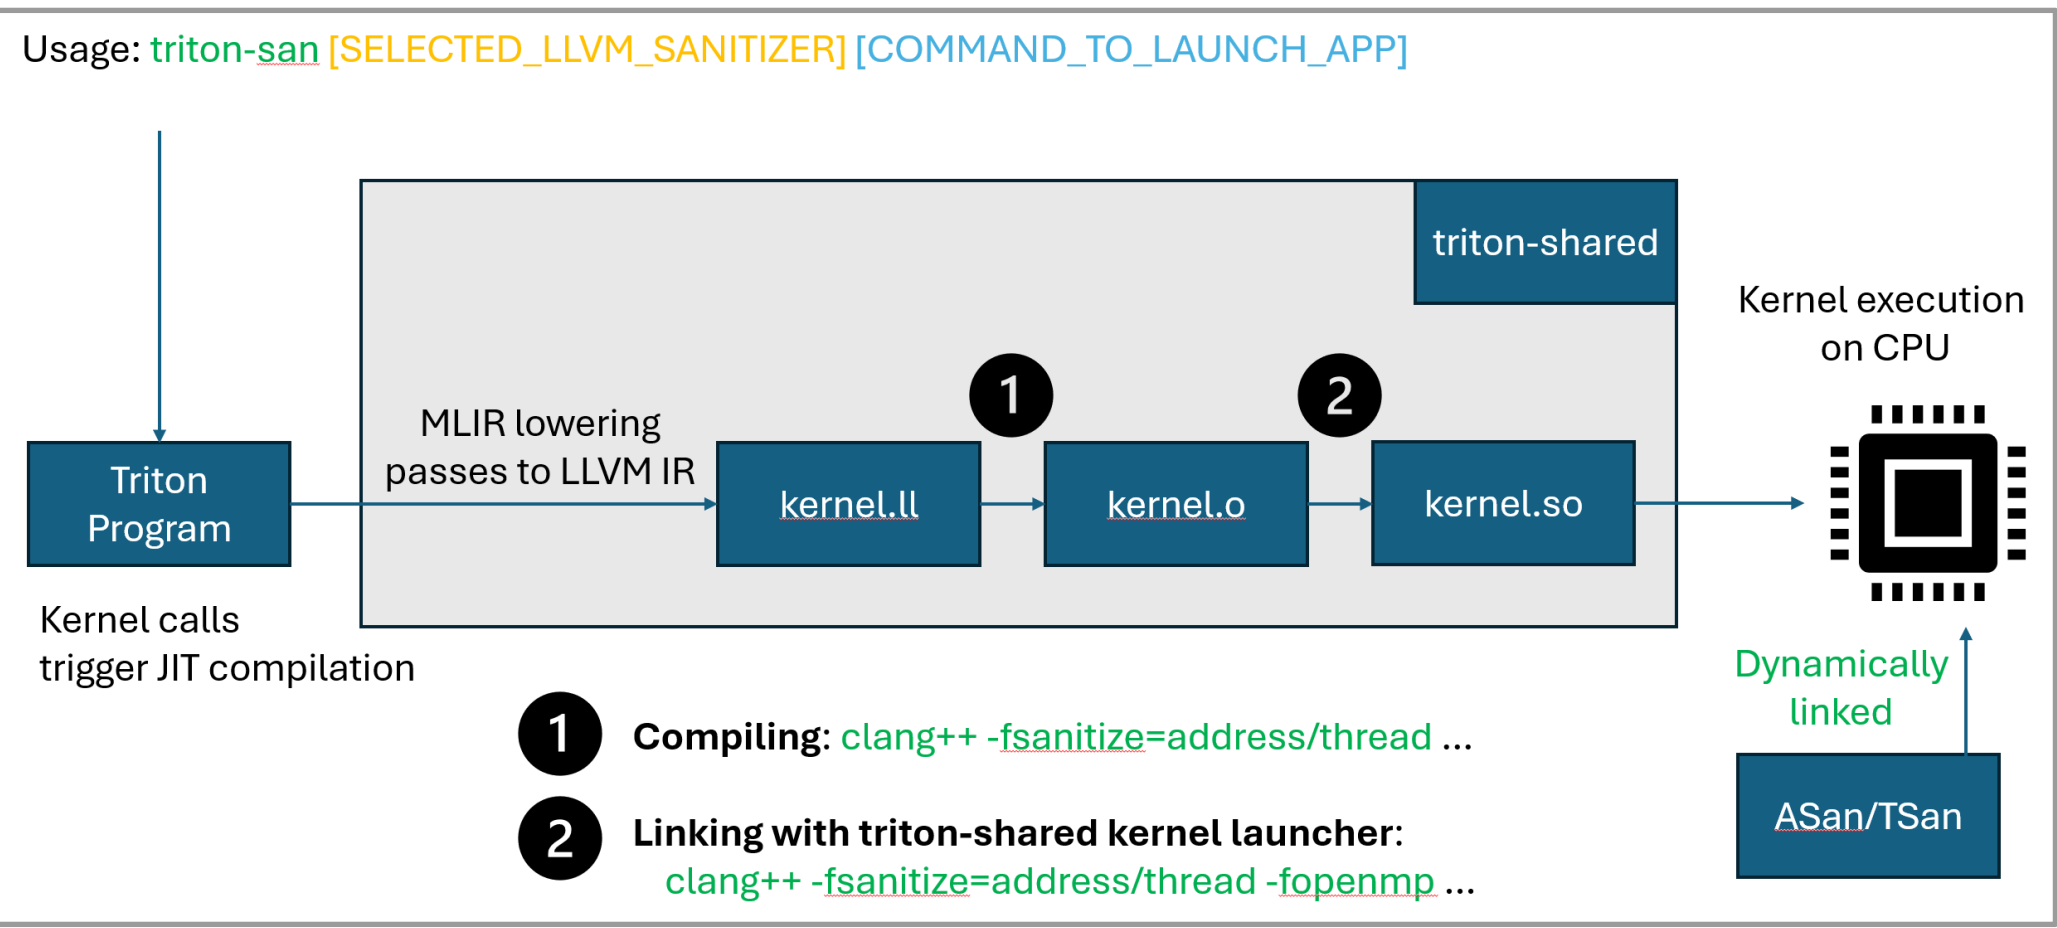
4. Triton-San Bug Report

```
WARNING: ThreadSanitizer: data race (pid=3823476)
  Write of size 8 at 0x726004bd800 by thread T44:
    #0 tsan_memcpy tsan_interceptors_memintrinsics.cpp:27:3
    #1 kernel data-race.py:14:35
    #2 _launch(int, int, int, void*, int, _object*) (.omp_outlined_debug_)
/tmp/tmpb3gw_jsv/main.cxx:26:11
    #3 _launch(int, int, int, void*, int, _object*) (.omp_outlined)
/tmp/tmpb3gw_jsv/main.cxx:20:5
    #4 __kmp_invoke_microtask <null> (libomp.so+0xc3d98)
    #5 _launch(int, int, int, void*, int, _object*) /tmp/tmpb3gw_jsv/main.cxx:20:5
(__triton_shared_ref_cpu_kernel_launcher.so+0x764a)

Previous write of size 8 at 0x726004bd800 by main thread:
    #0 tsan_memcpy tsan_interceptors_memintrinsics.cpp:27:3
    #1 kernel data-race.py:14:35
    #2 _launch(int, int, int, void*, int, _object*) (.omp_outlined_debug_)
/tmp/tmpb3gw_jsv/main.cxx:26:11
    #3 _launch(int, int, int, void*, int, _object*) (.omp_outlined)
/tmp/tmpb3gw_jsv/main.cxx:20:5
    #4 __kmp_invoke_microtask <null> (libomp.so+0xc3d98)
    #5 _launch(int, int, int, void*, int, _object*) /tmp/tmpb3gw_jsv/main.cxx:20:5
(__triton_shared_ref_cpu_kernel_launcher.so+0x764a)
    #6 <null> <null> (python3.12+0x581a8f)
```

5. System Overview of Triton-San

On-the-fly Dynamic Analysis for Triton Kernels




6. Effectiveness of Triton-San

- Developed micro-benchmarks to test the effectiveness of correctness tools when tackling Triton kernels
 - Each benchmark contains one injected bug
- | Error Type | Compute Sanitizer [6] | Triton-San |
|-----------------|-----------------------|------------|
| Buffer overflow | ✗ | ✓ |
| Data race | ✗ | ✓ |
- Triton-San detected all injected bugs with no false alarms
 - Compute Sanitizer showed several limitations
 - No data races detected (none occur in shared memory)
 - Detected buffer overflows only on very large tensors

7. Evaluations of Runtime Overhead

- Performed a preliminary performance evaluation
- Used four benchmarks from Triton tutorials
 - vector-add
 - fused-softmax
 - matrix-multiplication
 - low-memory-dropout
- Launched Triton kernels with 16 CPU threads as the baseline version



8. Other Correctness Tools for Triton

- For Triton kernels running on Nvidia GPUs, using correctness tools from the CUDA toolkit (i.e., *Nvidia Compute Sanitizer* [6]) can serve as an alternative for debugging and validation.
 - Racecheck: targeting only **shared-memory data races**
 - Memcheck: targeting memory issues (e.g, buffer overflow, memory leak)
 - Initcheck: targeting uninitialized memory accesses
- Our evaluations show that Compute Sanitizer is not very effective at identifying Triton bugs
- ConSan, a new tool for Triton, is under active development